

# AlphaFold Analysis

Emily Ignatoff (A16732102)

## AlphaFold Analysis of find-a-gene sequence:

Here we analyze our AlphaFold structure prediction models for our HIV sequence:

```
results_dir <- "hivpr_monomer_94b5b/"
```

First, I will read in each generated pdb file from the results directory:

```
pdb_files <- list.files(path=results_dir,  
                        pattern="*.pdb",  
                        full.names = TRUE)  
  
basename(pdb_files)
```

```
[1] "hivpr_monomer_94b5b_unrelaxed_rank_001_alphafold2_ptm_model_5_seed_000.pdb"  
[2] "hivpr_monomer_94b5b_unrelaxed_rank_002_alphafold2_ptm_model_4_seed_000.pdb"  
[3] "hivpr_monomer_94b5b_unrelaxed_rank_003_alphafold2_ptm_model_1_seed_000.pdb"  
[4] "hivpr_monomer_94b5b_unrelaxed_rank_004_alphafold2_ptm_model_3_seed_000.pdb"  
[5] "hivpr_monomer_94b5b_unrelaxed_rank_005_alphafold2_ptm_model_2_seed_000.pdb"
```

Now, I will superimpose each pdb on top of the others:

```
library(bio3d)
```

Warning: package 'bio3d' was built under R version 4.3.3

```
pdbs <- pdbaln(pdb_files, fit=TRUE, exefile="msa")
```

Reading PDB files:

```
hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_001_alphafold2_ptm_model_5_seed_000.p
hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_002_alphafold2_ptm_model_4_seed_000.p
hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_003_alphafold2_ptm_model_1_seed_000.p
hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_004_alphafold2_ptm_model_3_seed_000.p
hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_005_alphafold2_ptm_model_2_seed_000.p
.....
```

Extracting sequences

```
pdb/seq: 1   name: hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_001_alphafold2_ptm
pdb/seq: 2   name: hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_002_alphafold2_ptm
pdb/seq: 3   name: hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_003_alphafold2_ptm
pdb/seq: 4   name: hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_004_alphafold2_ptm
pdb/seq: 5   name: hivpr_monomer_94b5b/hivpr_monomer_94b5b_unrelaxed_rank_005_alphafold2_ptm
```

pdbs

```

1                                     .               .               .               .               50
[Truncated_Name:1]hivpr_mono  PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGI
[Truncated_Name:2]hivpr_mono  PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGI
[Truncated_Name:3]hivpr_mono  PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGI
[Truncated_Name:4]hivpr_mono  PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGI
[Truncated_Name:5]hivpr_mono  PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGI
*****
1                                     .               .               .               .               50

51                                     .               .               .               .               99
[Truncated_Name:1]hivpr_mono  GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:2]hivpr_mono  GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:3]hivpr_mono  GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:4]hivpr_mono  GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
[Truncated_Name:5]hivpr_mono  GGFIKVRQYDQILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNF
*****
51                                     .               .               .               .               99
```

Call:

```
pdbaln(files = pdb_files, fit = TRUE, exefile = "msa")
```

Class:

```
pdbs, fasta
```

Alignment dimensions:

5 sequence rows; 99 position columns (99 non-gap, 0 gap)

+ attr: xyz, resno, b, chain, id, ali, resid, sse, call

Now, I will generate the structural difference between these coordinate sets using `rmsd()`:

```
rd <- rmsd(pdbbs, fit=T)
```

Warning in `rmsd(pdbbs, fit = T)`: No indices provided, using the 99 non NA positions

```
range(rd)
```


```
[1] 0.000 1.206
```

Heatmap of RMSD values:

```
library(pheatmap)
```

Warning: package 'pheatmap' was built under R version 4.3.3

```
colnames(rd) <- paste0("m",1:5)
rownames(rd) <- paste0("m",1:5)
pheatmap(rd)
```



I will be using a human HSP70 protein from the PDB database as my reference protein for the following steps:

```
pdb <- read.pdb("1hsg")
```


Note: Accessing on-line PDB file

Now, I will plot the pLDDT values of each model against each other:

```
plotb3(pdb$b[1,], typ="l", lwd=2, sse=pdb)
```

Warning in plotb3(pdb\$b[1, ], typ = "l", lwd = 2, sse = pdb): Length of input 'sse' does not equal the length of input 'x'; Ignoring 'sse'

```
points(pdb$b[2,], typ="l", col="red")
points(pdb$b[3,], typ="l", col="blue")
points(pdb$b[4,], typ="l", col="darkgreen")
points(pdb$b[5,], typ="l", col="orange")
abline(v=100, col="gray")
```



I will now find the the most “rigid” or consistent portion across models using the `core.find()` function:

```
core <- core.find(pdbbs)
```

```
core size 98 of 99  vol = 3.66
core size 97 of 99  vol = 2.756
core size 96 of 99  vol = 2.236
core size 95 of 99  vol = 1.751
core size 94 of 99  vol = 1.386
core size 93 of 99  vol = 0.991
core size 92 of 99  vol = 0.769
core size 91 of 99  vol = 0.568
core size 90 of 99  vol = 0.422
FINISHED: Min vol ( 0.5 ) reached
```

```
core.inds <- print(core, vol=0.5)
```

```
# 91 positions (cumulative volume <= 0.5 Angstrom^3)
  start end length
1     3   3      1
2     7  96     90
```

Writing the superimposed coordinates to a new file for use in Mol\*

```
xyz <- pdbfit(pdb, core.inds, outpath="corefit_structures")
```




Figure 1: Uncertainty/disorder of structures in Mol\*

I will now use RMSF to observe the conformational variance across the structure:

```
rf <- rmsf(xyz)  
plotb3(rf, sse=pdb)
```

Warning in plotb3(rf, sse = pdb): Length of input 'sse' does not equal the length of input 'x'; Ignoring 'sse'

```
abline(v=100, col="gray", ylab="RMSF")
```



Note: I am unsure why the sse function is not working using an existing PDB file as my reference protein :(

## Predicted Alignment Error for Domains

```
library(jsonlite)
```

Warning: package 'jsonlite' was built under R version 4.3.3

```
pae_files <- list.files(path=results_dir,  
                        pattern=".*model.*\\.json",  
                        full.names = TRUE)
```

```
pae1 <- read_json(pae_files[1],simplifyVector = TRUE)
pae2 <- read_json(pae_files[2],simplifyVector = TRUE)
pae3 <- read_json(pae_files[3],simplifyVector = TRUE)
pae4 <- read_json(pae_files[4],simplifyVector = TRUE)
pae5 <- read_json(pae_files[5],simplifyVector = TRUE)

attributes(pae1)
```

```
$names
[1] "plddt" "max_pae" "pae" "ptm"
```

```
head(pae1$plddt)
```

```
[1] 77.56 88.25 90.88 92.88 90.56 94.12
```

Which model has the best (lowest) max PAE?

```
pae1$max_pae
```

```
[1] 17.67188
```

```
pae2$max_pae
```

```
[1] 16.9375
```

```
pae3$max_pae
```

```
[1] 19.25
```

```
pae4$max_pae
```

```
[1] 21.48438
```

```
pae5$max_pae
```


```
[1] 20.95312
```




PAE 2 is lowest in this case!

I will plot each of these PAE in the same range:

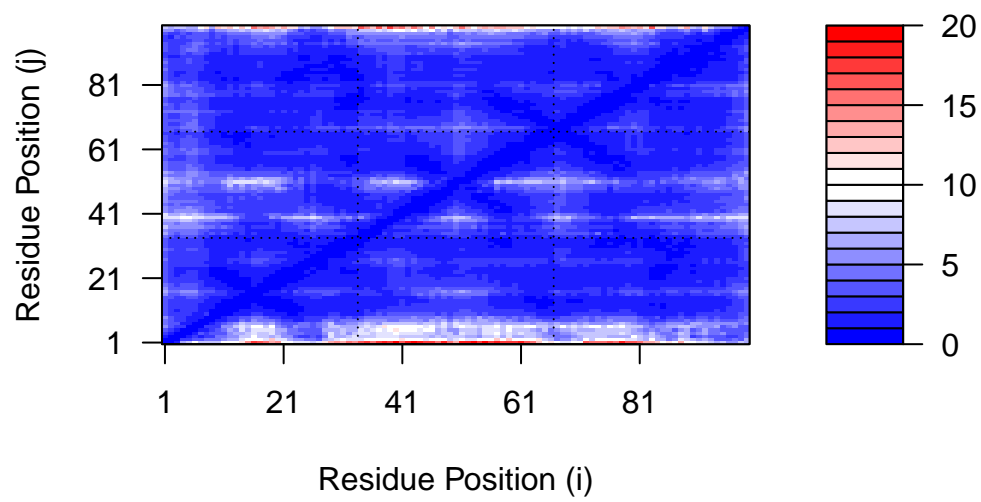
```
plot.dmat(pae1$paes,  
          xlab="Residue Position (i)",  
          ylab="Residue Position (j)",  
          grid.col = "black",  
          zlim= c(0,20))
```




```
plot.dmat(pae2$paes,  
          xlab="Residue Position (i)",  
          ylab="Residue Position (j)",  
          grid.col = "black",  
          zlim= c(0,20))
```




```
plot.dmat(pae3$pae,  
  xlab="Residue Position (i)",  
  ylab="Residue Position (j)",  
  grid.col = "black")
```



```
plot.dmat(pae4$pae,  
          xlab="Residue Position (i)",  
          ylab="Residue Position (j)",  
          grid.col = "black")
```



```
plot.dmat(pae5$pae,  
  xlab="Residue Position (i)",  
  ylab="Residue Position (j)",  
  grid.col = "black")
```



## Residue Conservation from Alignment File

Generate alignment file:

```
aln_file <- list.files(path=results_dir,
                      pattern=".a3m$",
                      full.names = TRUE)
aln_file
```

```
[1] "hivpr_monomer_94b5b/hivpr_monomer_94b5b.a3m"
```

```
aln <- read.fasta(aln_file[1], to.upper = TRUE)
```

```
[1] " ** Duplicated sequence id's: 101 **"
```

How many sequences in this alignment?


```
dim(aln$ali)
```

```
[1] 5378 132
```

I will score residue conservation using the `conserv()` function:

```
sim <- conserv(aln)
```

```
plotb3(sim[1:99], sse=trim.pdb(pdb, chain="A"),  
        ylab="Conservation Score")
```



Let's look at the sequences which will stand out in a consensus sequence

```
con <- consensus(aln, cutoff = 0.9)  
con$seq
```

```
[1] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[19] "-" "-" "-" "-" "-" "-" "D" "T" "G" "A" "-" "-" "-" "-" "-" "-" "-" "-"  
[37] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[55] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[73] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[91] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[109] "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-" "-"  
[127] "-" "-" "-" "-" "-" "-"
```

```
m1.pdb <- read.pdb(pdb_files[1])  
occ <- vec2resno(sim[1:length(unique(m1.pdb$atom$resno))], m1.pdb$atom$resno)  
write.pdb(m1.pdb, o=occ, file="m1_conserv.pdb")
```



Figure 2: Color by Occupancy in Mol\*